

# End-to-End PPO-Based Reinforcement Learning Framework for Scalable 0/1 Knapsack Problem Solving

From Data Generation to Large-Scale Generalization

Gang Lin  
Student ID: 2874886


**University of Birmingham**

19/08/2025

# The 0/1 Knapsack Problem: Definition & Context



(a) The Knapsack Problem Analogy.



(b) Computational Complexity Classes.

## The 0/1 Knapsack Problem

Given  $n$  items with weights and values, choose a subset.

- **Objective:** Maximize the total value.
- **Constraint:** The total weight must not exceed the knapsack's capacity.
- **Rule:** Each item is either taken (1) or left (0), no fractions

## Applications & Complexity

- **Real-world Applications:** Portfolio selection, resource allocation, logistics.
- **Problem Variants:** Many variants exist by altering constraints or item properties. The 0/1 version is the most common.
- **Computational Hardness:** KP is **NP-complete**. There is no known polynomial-time exact solution.

## 1. Exact and Approximate Algorithms

Algorithm	TC.	SC.	Limitations
Dynamic Programming	$O(n \cdot C)$	$O(n \cdot C)$	High memory usage; infeasible for large capacity $C$
Branch & Bound	$O(2^n)$ (worst-case)	$O(n^2)$	Exponential runtime in worst case; performance depends on bounding quality
Greedy	$O(n \log n)$	$O(n)$	No performance guarantee; poor approximation in worst case
Genetic Algorithm	$O(G \cdot P \cdot n)$	$O(P \cdot n)$	Parameter-sensitive; may converge prematurely
Simulated Annealing	$O(\text{iter} \cdot n)$	$O(n)$	Slow convergence; sensitive to cooling schedule

## 2. Modern Neural Network Approaches (for Knapsack Problem)

Work	Type & Gen.	Architecture / Algorithm	Key Results (Accuracy / Speed)
Bello (2017)	C / Fixed	PtrNet / REINFORCE	Near-optimal on small scale; Faster than exact algorithms
Yildiz (2022)	C / Fixed	Transformer / DQN	~92.7% accuracy; ~40x faster than DP
Abid (2023)	C / Fixed	MLP / SL	Near-optimal but slower inference than heuristics
Cappart (2021)	I / Yes	DRL + CP Solver	Proves optimality; Outperforms standalone RL/CP

**Note:** TC. = Time Complexity, SC. = Space Complexity.  $n$ : number of items,  $C$ : knapsack capacity,  $G$ : generations,  $P$ : population size, iter: iterations. **C** = Constructive; **I** = Improvement; **Gen.** = Generalization. PtrNet = Pointer Network; MLP = Multi-Layer Perceptron; SL = Supervised Learning; CP = Constraint Programming.

**My Contribution:** A generalizable RL framework that solves knapsack problems of arbitrary size (train on  $N$ , test on  $> N$ ) with 70% accuracy using PPO, enabling scalable and robust decision-making.

# From Dynamic Programming (DP) to Reinforcement Learning (RL)

## 1. Dynamic Programming Recurrence for KP

State value is the maximum value by:

- **Skip item  $i$ :** Value is the future reward from the remaining state,  $V(i-1, w)$ .
- **Take item  $i$ :** Value is the immediate reward  $v_i$  + future reward from the new state,  $V(i-1, w - w_i)$ .

$$V(i, w) = \begin{cases} V(i-1, w) & \text{if } w_i > w \\ \max(V(i-1, w), v_i + V(i-1, w - w_i)) & \text{if } w_i \leq w \end{cases}$$

## 2. The Bellman Equation (Element-wise)

State value is the expected immediate reward plus the discounted expected future reward.

$$V_{\pi}(s) = \sum_{a \in A} \pi(a|s) \left[ \sum_{r \in R} p(r|s, a)r + \gamma \sum_{s' \in S} p(s'|s, a)V_{\pi}(s') \right]$$

- **DP** recurrence is a specific, deterministic instance of the Bellman equation.

## Modeling KP as an RL Problem

- **State ( $s_t$ ):** The set of available items and the current remaining knapsack capacity.
- **Action ( $a_t$ ):** The selection of one item from the available set.
- **Reward ( $R_{t+1}$ ):** The value ( $v_i$ ) of the selected item.
- **Policy ( $\pi_{\theta}(a|s)$ ):** A neural network mapping states to action probabilities.

# From Bellman to PPO: An Evolutionary Path

## 1 Bellman Equation

$$V^\pi(s) = \mathbb{E}[r + \gamma V^\pi(s')]$$

## 2 Value Function Approx.

$$V_\phi(s) \approx V^\pi(s)$$

## 3 REINFORCE (Policy Gradient)

$$\nabla_\theta J = \mathbb{E}[\nabla_\theta \log \pi_\theta(a|s) \cdot G_t]$$

## 4 A2C (Actor-Critic)

$$\nabla_\theta J \propto \sum_t \nabla_\theta \log \pi_\theta(a_t|s_t) \cdot \hat{A}_t$$

$$(\hat{A}_t = r + \gamma V_\phi(s') - V_\phi(s): \text{TD-based})$$

## 5 PPO (Clipped Objective)

$$\mathcal{L}^{\text{CLIP}}(\theta) = \mathbb{E} \left[ \min \left( r_t \hat{A}_t, \text{clip}(r_t, 1 \pm \epsilon) \hat{A}_t \right) \right]$$

$$\text{where } r_t = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$$

## The Evolutionary Path to PPO

**Bellman** → **Value Approx.** Replace exact  $V^\pi$  with a learnable  $V_\phi(s)$  (e.g., neural net). Enables scalability to large or continuous state spaces.

**Value Approx.** → **REINFORCE** Shift from value-based to direct policy optimization. More suitable for stochastic or complex action spaces.

**REINFORCE** → **A2C** Replace Monte Carlo return  $G_t$  with TD-based advantage  $\hat{A}_t$ . Reduces variance, enables online updates, and improves sample efficiency.

**A2C** → **PPO** Replace policy gradient with clipped surrogate objective. Prevents destructive updates and stabilizes training.

# PPO: Algorithm and Architecture

## Training Algorithm




Figure: The PPO training loop using an Actor-Critic framework.

- Multiple optimization on the same minibatch.

## Model Architecture

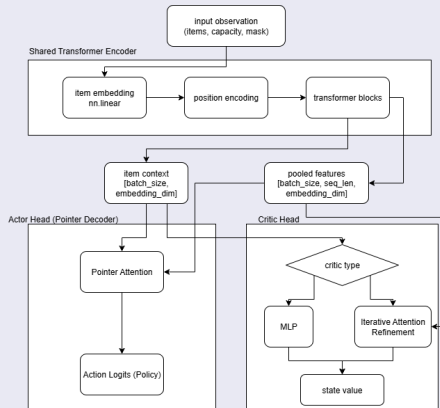


Figure: The model has two heads: one for the policy (Actor) and one for the value (Critic).

- Actor and Critic share the same encoder.

## Dataset Specification

We generated three distinct datasets for training, validation, and testing to ensure a robust evaluation of the model's generalization capabilities.

Parameter	Training Set	Validation Set	Test Set
Item Count Range ( $n$ )	5 to 50	5 to 50	<b>5 to 200</b>
Step Size	5	5	5
Instances per Size	100	30	50
<b>Total Instances</b>	<b>1,000</b>	<b>300</b>	<b>1,950</b>

## Item Properties

- Weights ( $w_i$ ) and values ( $v_i$ ) are integers sampled uniformly from  $U[1, 100]$ .
- There is no correlation between an item's weight and its value.
- All inputs are normalized before being fed to the model.

## Problem Instance Constraints

- The knapsack capacity ( $C$ ) is set relative to the total weight of all items ( $\sum w_i$ ).
- The ratio  $\frac{C}{\sum w_i}$  is randomly sampled from  $U[0.1, 0.9]$ .


# Results: Accuracy and Inference Time



(a) Mean Relative Error (MRE) vs. Problem Size.

## Key Findings: Accuracy

- Our PPO model maintains a low Mean Relative Error (MRE), demonstrating high solution quality and strong generalization.
- Pointer Network shows a higher error rate.
- The pure MLP model fails to generalize effectively.



(b) Inference Time vs. Problem Size.

## Key Findings: Inference Time

- PPO's inference time is practical for large instances.
- Pointer Network is faster but less accurate.
- MLP is the fastest but provides poor solutions.



## Performance Summary

### ■ PPO vs. Pointer Network (Accuracy):

- **Algorithmic Superiority:** PPO's Actor-Critic (TD) method provides low-variance updates.
- **Architectural Advantage:** The **Transformer** encoder captures the global, combinatorial nature of the problem more effectively than a sequential **LSTM**.
- **Framework Robustness:** Leveraging **Stable Baselines 3** provides key stabilizations like adaptive observation normalization ('VecNormalize').

### ■ PPO vs. Pointer Network (Speed):

- **Core Architecture:** **Transformer** is more computationally intensive than **LSTM**.
- **Model Components:** Extra Critic Network requires extra computation.
- **Evaluation Method:** `Stable_baseline3` cannot support batch evaluation.

## Effective Training Techniques

The success of the framework relies on several key techniques:

### ■ Input Normalization:

Normalizing item attributes ( $w_i, v_i$ ) and the knapsack capacity ( $C$ ) is crucial.

### ■ Observation & Reward Normalization:

Using 'VecNormalize' for both observations and rewards stabilizes the learning process significantly.

### ■ Heuristic Preprocessing:

Sorting items by value-density ( $v_i/w_i$ ) before feeding them to the model provides a strong inductive bias and improves performance.

## Architectural Exploration

- **The "Simple Critic" Anomaly:**

A simple MLP Critic achieved higher accuracy (70%) than a more complex attention-based head (60%). Future work should investigate if this is due to optimization challenges or a regularization effect.

- **Global State Representation ('[CLS]' Token):**

Initial experiments with a '[CLS]' token for global state representation surprisingly decreased performance. This warrants further investigation.

- **Hyperparameter Tuning:**

While a 3-layer MLP Critic works well, its optimal width and the interplay with network depth remain open questions for further tuning.

## Problem Formulation & Reward Shaping

- **Explore Alternative Formulation:** Our model uses a **"Decision" formulation** (select one from all remaining items). An alternative **"Selection" formulation** (decide 'take' or 'skip' for items sequentially) could be investigated.

- **Advanced Reward Shaping:** For the current "Decision" model, an initial attempt at adding a final shaping reward (to encourage a fuller knapsack) decreased accuracy. Further research into more advanced shaping techniques (e.g., potential-based rewards) is needed.